

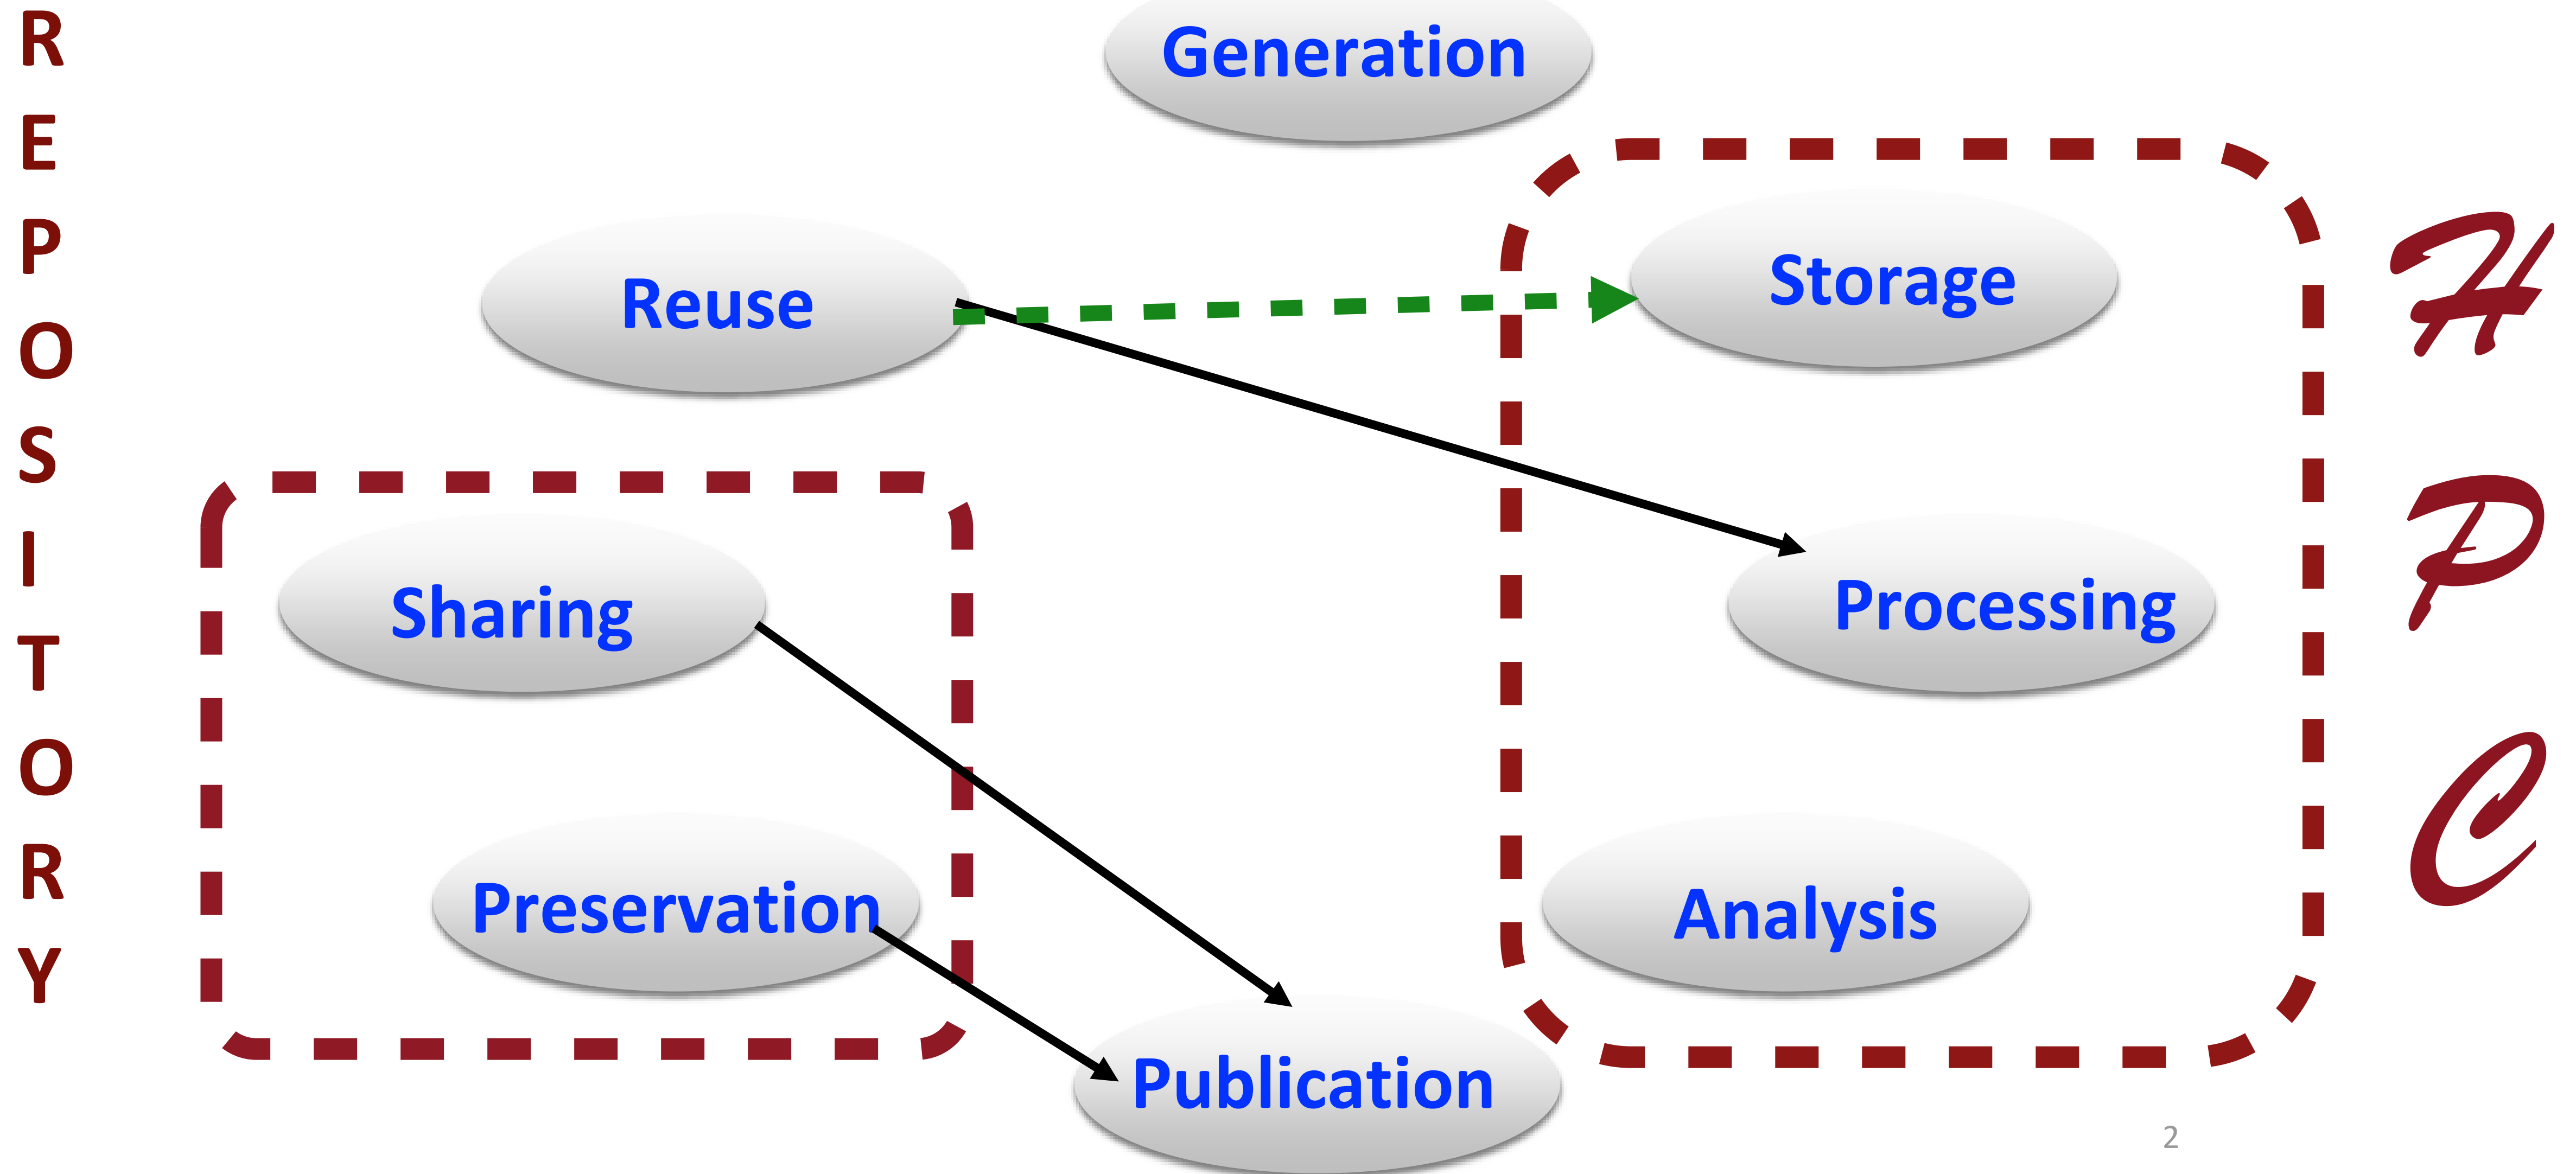
Maryland Advanced Research Computing Center

HHPC@JHU



Jaime E. Combariza, PhD
Director

Data Lifecycle



Collecting Big Data




High Res.
telescopes




Social Networks



Virtual
Libraries



Brain Science



Predictive
Computational
Science



Satellites



Human
Genomics



Smart
Networks

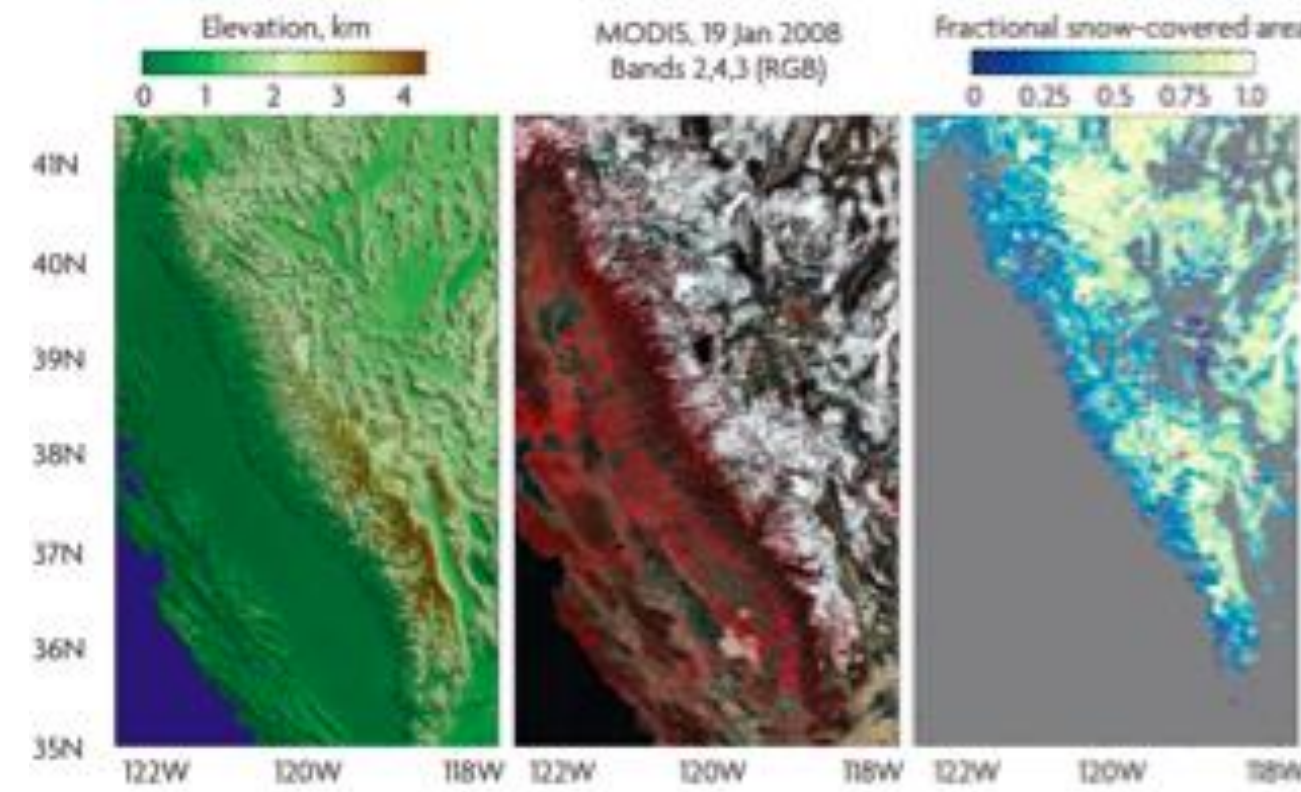


Financial
Markets



Large Scale
Experiments

Data-Rich Environment (Effects)



Eco-Informatics



Biomedical Informatics



Digital Humanities



Astro-Informatics




Data Driven Applications

Opportunities



Innovation, Competition and Productivity

What's being done


- WH 2012 200M R&D for big data
- NSF & NIH (many programs)
- Gordon and Betty Moore Foundation
- Many schools VIDA, DI², 
- States, e.g.: Mass
- Xsede (9/2/14)



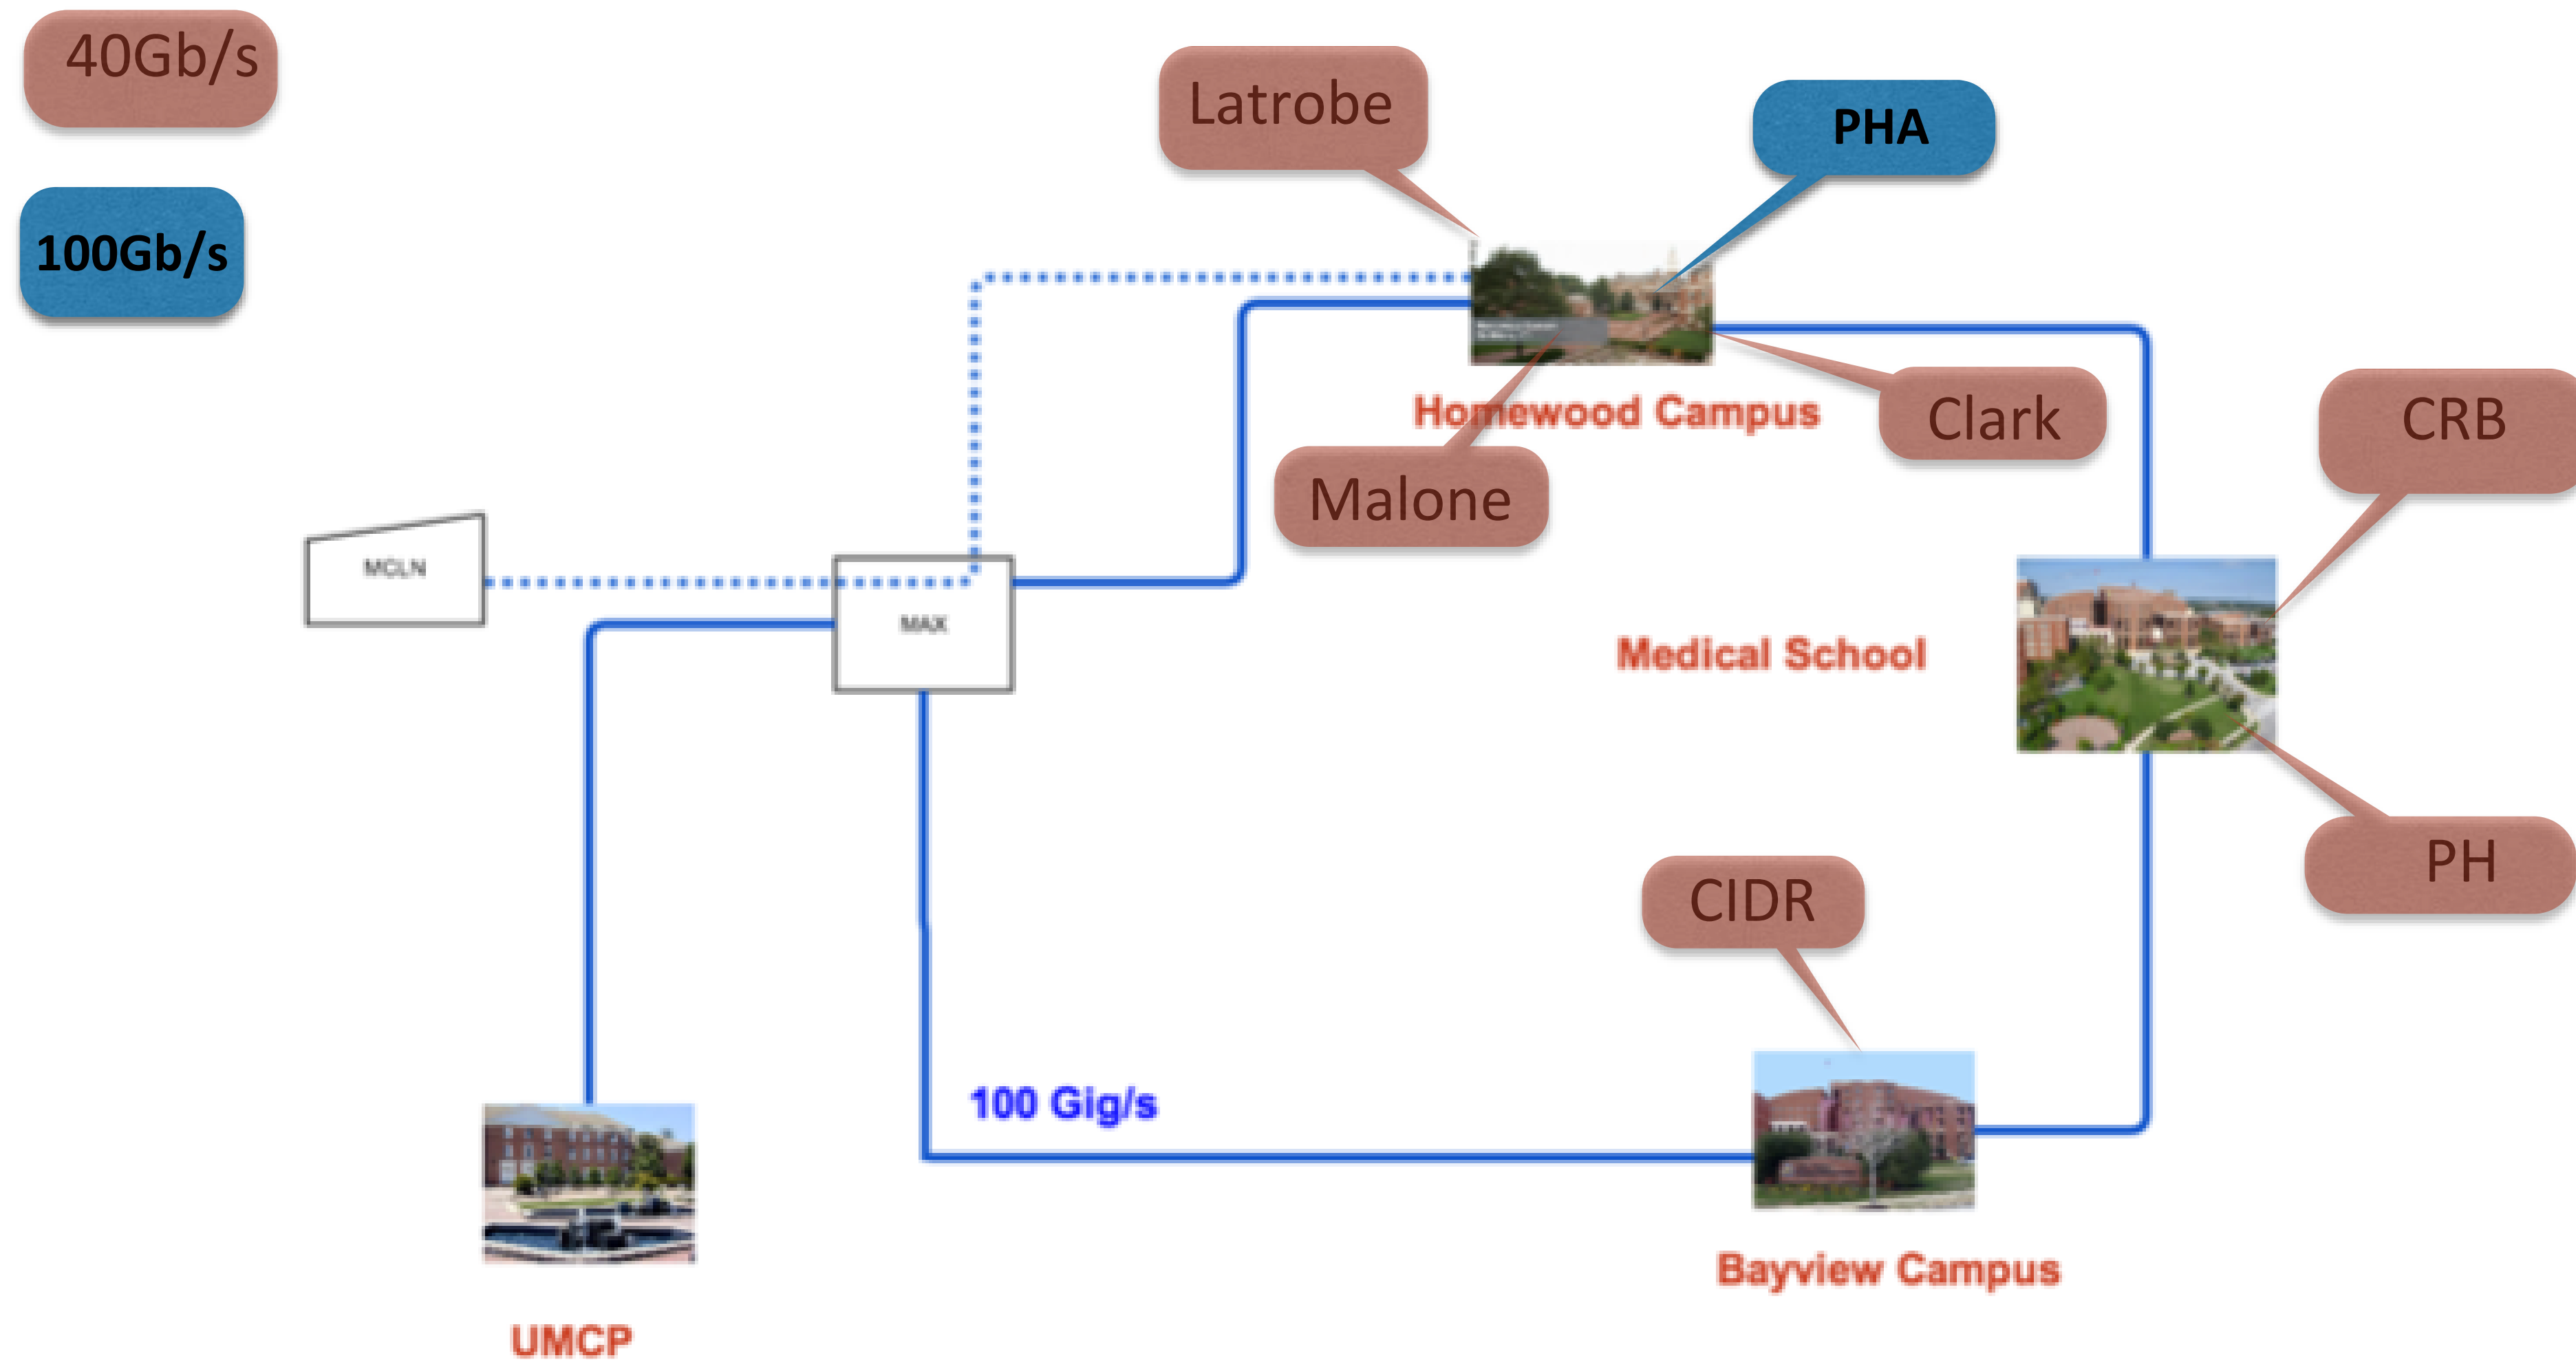
What's Maryland doing?



Data center and data-rich environment




Science DMZ @ JHU



HoRNET

HPC Resources & Model

- Approx 18,000 cores and 20 Petabytes storage



KSAS: 10.5M Quarter

WSE: 10.5M/Q

SOM: 6.2M/Q

PH: 5.8M/Q

UMCP: 5.8M/Q

Hardware

Count	Description
648	Regular compute nodes, 128 GB RAM, Haswell 2.5GHz, 24 cores
50	Large memory nodes, 1 Tbyte RAM, ivy Bridge 3.0GHz, 48 cores
24	GPU nodes, 48 Nvidia K40 GPUs, 24 cores per node
2 PB	High Performance File System (CIFS)
14 PB	zFS File System
FDR-14 Infiniband Connectivity 2:1 blocking	
$R_{\text{peak}} = 771 \text{ TFLOPs}$, $R_{\text{Max}} = 617 \text{ TFLOPs}$, (61st place on Top500)	



Governance



- Facility Governing Board (3 members JHU, 3 members UMCP)
- Scientific Management Committee (4 faculty members from JHU and UMCP).
 - Policies
 - Allocation models
 - Utilization

Timeline

- Data center: End of November
- Network install “mid December”
- Hardware deployment Dec 1 - Jan 2015
- Production system:



Co-Location

- Priority
- Early Spring we will have better guidelines and policies
- Cluster/server hosting and management
- Cluster/server hosting
- Virtual servers



Next . . .

- Website
 - FAQs
 - Instruction on how to connect
 - Request accounts
 - Queueing system and queues
- Software & Application support
- Add Condos
- Documentation



Summary

- Create state of the art HPC facility (current and future needs)
- Develop an HPC ecosystem:
 - To foster collaboration and advance research agendas
 - Develop or improve tools for more effective data analysis
 - Resource to attract and retain top researchers
 - Facility to properly house HPC resources with different requirements (Co-location)
 - Researchers concentrate on science not on IT



Information

- combariza@jhu.edu
- Web site (soon) marcc.jhu.edu









